

# CMP 713 Data Mining - Mining Titanic Dataset

Erol ÖZKAN

Hacettepe University - Computer Engineering  
erolozkan@outlook.com

18 Mayıs 2020

## Abstract

In this work, we aim to predict what kind of people are likely to survive in the Titanic dataset. In short, we use several data mining methods to predict the survived passengers in the testing dataset. We do explanatory data analysis and create some illustrative data visualization plots. Then, we do feature engineering and missing value imputation to create more meaningful variables. Finally, we build RandomForrest, SVM and GBM models to predict the Survival variable in the testing datasets. We compare these there results.

## 1 INTRODUCTION

The remainder of the document is organized as follows. In Section 2, we explain the Titanic dataset. In Section 3, we explain the EDA. Missing value imputation is explained in Section 5. In Section 4 and Section 6 we show feature engineering steps. Predictive models are explained in the Section 7. Results are shown in Section 8. Finally, we conclude our work in Section 8.

## 2 DATASET

Titanic dataset is used for this project. It contains the survival information of Titanic ship which sank after colliding with an iceberg on April 15, 1912. It killed 1502 out of 2224 including both all the passengers and its crew. One of the reasons this tragedy led to such loss of life was that there were not enough lifeboats for everyone on the ship. This tragedy resulted in better safety regulations for ships.

## 3 FIRST STEPS

### 3.1 Import Required Packages

```
library('ggplot2')  
library('ggthemes')  
library('scales')  
library('mice')  
library('knitr')  
library('dplyr')  
library('caret')  
library('randomForest')
```

```
library('gridExtra')
library('corrplot')
#library('Hmisc')
#library('ROCR')
```

## 3.2 Read Data

```
train <- read.csv('dataset/train.csv', stringsAsFactors = F, na.strings = c("NA", "")) # Read train data
test  <- read.csv('dataset/test.csv', stringsAsFactors = F, na.strings = c("NA", "")) # Read test data
full  <- bind_rows(train, test) # Concatenate train and test dataset
```

# 4 EXPLORATORY DATA ANALYSIS

## 4.1 Check Data

Let's check the dataset first.

```
str(full)
```

```
## 'data.frame': 1309 obs. of 12 variables:
## $ PassengerId: int 1 2 3 4 5 6 7 8 9 10 ...
## $ Survived : int 0 1 1 1 0 0 0 0 1 1 ...
## $ Pclass : int 3 1 3 1 3 3 1 3 3 2 ...
## $ Name : chr "Braund, Mr. Owen Harris" "Cumings, Mrs. John Bradley (Florence Briggs Thayer)"
## $ Sex : chr "male" "female" "female" "female" ...
## $ Age : num 22 38 26 35 35 NA 54 2 27 14 ...
## $ SibSp : int 1 1 0 1 0 0 0 3 0 1 ...
## $ Parch : int 0 0 0 0 0 0 0 1 2 0 ...
## $ Ticket : chr "A/5 21171" "PC 17599" "STON/O2. 3101282" "113803" ...
## $ Fare : num 7.25 71.28 7.92 53.1 8.05 ...
## $ Cabin : chr NA "C85" NA "C123" ...
## $ Embarked : chr "S" "C" "S" "S" ...
```

There are 1309 observations in total. The training dataset has 891, the test dataset has 418 observations. Note that “Survived” variable is missing from the testing dataset.

## 4.2 Describe Variables

Description of the variables are shown below.

Variable Name	Description
Survived	Survived (1) or died (0)
<u>Pclass</u>	Passenger's class
Name	Passenger's name
Sex	Passenger's sex
Age	Passenger's age
<u>SibSp</u>	Number of siblings/spouses aboard
Parch	Number of parents/children aboard
Ticket	Ticket number
Fare	Fare
Cabin	Cabin
Embarked	Port of embarkation

Figure 1: Description of the Dataset

### 4.3 Investigate the Data

Let's investigate the variables that contain missing values.

```
sapply(full, function(x) {sum(is.na(x))})
```

```
## PassengerId  Survived  Pclass     Name     Sex     Age
##           0      418         0         0         0      263
##      SibSp    Parch    Ticket   Fare    Cabin Embarked
##           0         0         0         1      1014         2
```

418 observations is missing the Survived column in the test dataset. Cabin is sparsely populated. Age is missing in some observations. In addition, Embarked is missing two values and Fare one value.

### 4.4 Explore the Data

Let's convert Sex, Pclass, Survived and Embarked variables into factors.

```
full$Sex <- as.factor(full$Sex)
full$Survived <- as.factor(full$Survived)
full$Pclass <- as.factor(full$Pclass)
full$Embarked <- as.factor(full$Embarked)
```

#### 4.4.1 Survived

Survived variable describes the survived people.

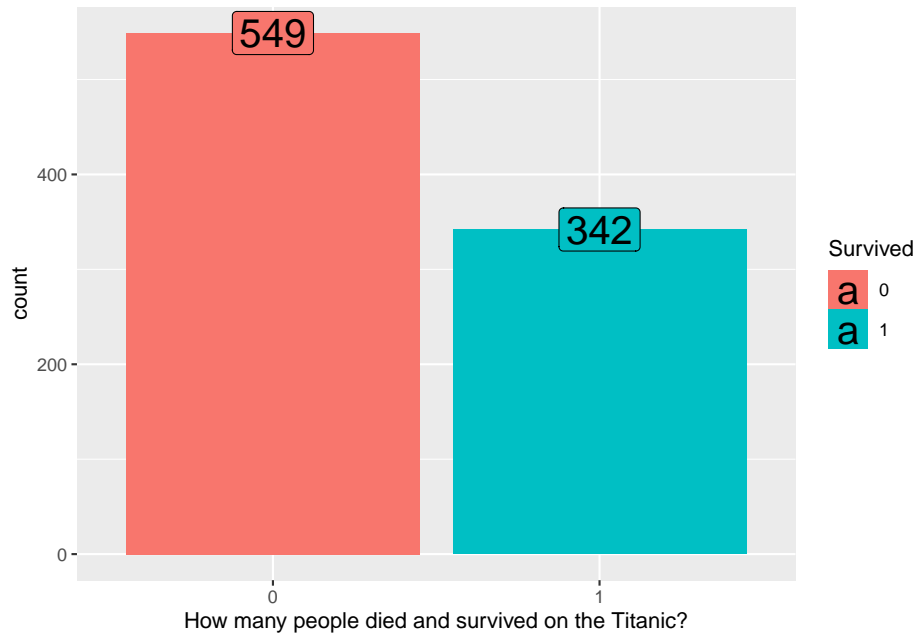


Figure 2: Survival of People

It is clear that 549 people died and 342 people survived in the training dataset.

#### 4.4.2 Sex

Let's visualize the Sex variable.

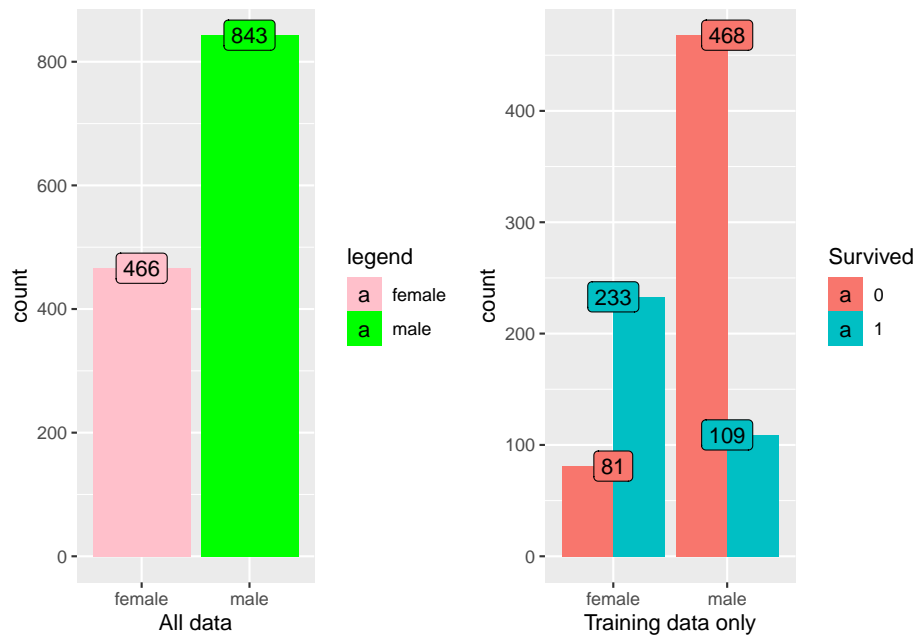


Figure 3: Survival rate of Male/Female Passangers

First plot shows that there are 466 female, 843 male passengers on the training and testing data. We can infer from second plot that 81.1% of the men died, only 25.8% of the women died in the training dataset.

### 4.4.3 PClass

Let's visualize the PClass variable.

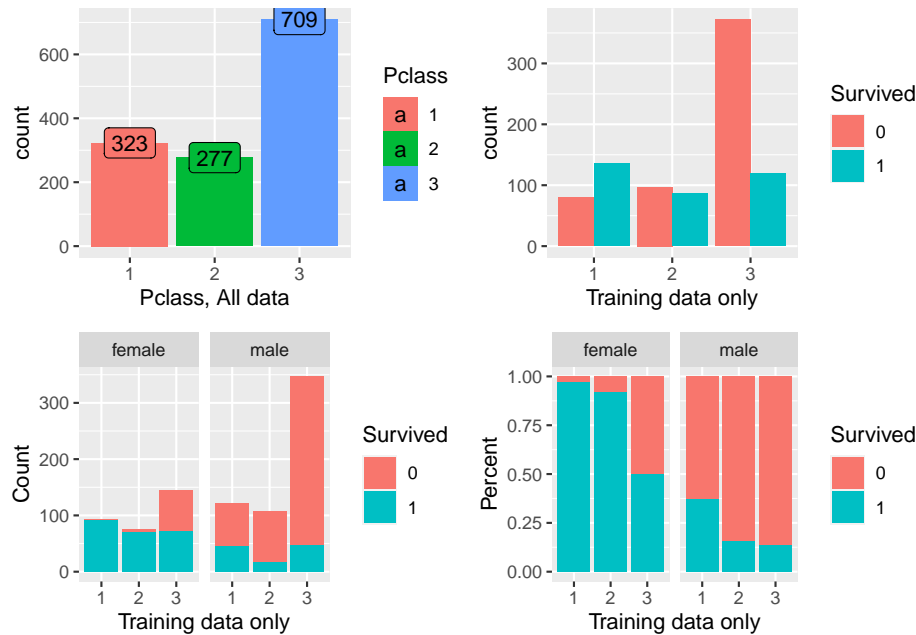


Figure 4: Passenger Class Information

First plot states that most people was in the 3rd class. According to second plot, majority of 1st and 2nd class passengers survived, while majority of 3rd class died. It is noticeable in 3rd and 4th charts that almost all women in 1st and 2nd class survived.

## 5 FEATURE ENGINEERING 1

### 5.1 Title

We can break the Name variable into additional meaningful variables such as; title, name and surname.

```
full$Title <- gsub('(.*, )|(\\.*)', '', full$Name) # Create Title variable
table(full$Sex, full$Title) # Show Title counts
```

```
##
##          Capt Col Don Dona  Dr Jonkheer Lady Major Master Miss Mlle Mme  Mr Mrs
## female    0  0  0  1  1      0  1  0  0 260  2  1  0 197
## male      1  4  1  0  7      1  0  2  61  0  0  0 757  0
##
##          Ms Rev Sir the Countess
## female    2  0  0      1
## male      0  8  1      0
```

Let's replace low title counts.

```
rare_title <- c('Dona', 'Lady', 'the Countess','Capt', 'Col', 'Don', 'Dr', 'Major', 'Rev', 'Sir', 'Jon')
full$Title[full$Title == 'Mlle']      <- 'Miss'
full$Title[full$Title == 'Ms']        <- 'Miss'
full$Title[full$Title == 'Mme']       <- 'Mrs'
full$Title[full$Title %in% rare_title] <- 'Rare Title'
table(full$Sex, full$Title)
```

```
##
##      Master Miss  Mr Mrs Rare Title
##  female      0 264   0 198      4
##  male       61   0 757   0     25
```

Let's visualize the Title variable.




Figure 5: Visualization of Title Variable

It is clear that Title is highly correlated with Survival information.

## 5.2 Family Size

Family size variable is based on number of siblings/spouse(s) and children/parents. We need to add 1 to include the passenger himself/herself.

```
full$Fsize <- full$SibSp + full$Parch + 1
```

To better understand the family size variable, let's use ggplot to visualize the family size and the survival.

```
ggplot(full[1:891,], aes(x = Fsize, fill = factor(Survived))) +
  geom_bar(stat='Count', position='dodge') +
  scale_x_continuous(breaks=c(1:11)) +
  labs(x = 'Family Size')
```





Figure 6: Family Size and Survival

It is clear that there is a survival penalty to single passengers and family sizes above 4. We can further make Family variable a categorical variable.

```
full$FsizeD[full$Fsize == 1] <- 'singleton'
full$FsizeD[full$Fsize < 5 & full$Fsize > 1] <- 'small'
full$FsizeD[full$Fsize > 4] <- 'large'
mosaicplot(table(full$FsizeD, full$Survived), main='Family Size - Survival', shade=TRUE)
```



This plot shows that the survival rate of large families is worse than singletons and small families.

### 5.3 Cabin

Let's take a look into the Cabin variable.

```
full$Cabin[1:28]
```

```
## [1] NA          "C85"        NA           "C123"       NA
## [6] NA          "E46"        NA           NA           NA
## [11] "G6"        "C103"       NA           NA           NA
## [16] NA          NA           NA           NA           NA
## [21] NA          "D56"       NA           "A6"         NA
## [26] NA          NA           "C23 C25 C27"
```

As you can see, there are lots of missing values. But, we can also infer Deck variable from Cabin variable.

```
strsplit(full$Cabin[2], NULL)[[1]]
```

```
## [1] "C" "8" "5"
```

```
full$Deck<-factor(sapply(full$Cabin, function(x) strsplit(x, NULL)[[1]][1]))
full$Deck[1:28]
```

```
## [1] <NA> C    <NA> C    <NA> <NA> E    <NA> <NA> <NA> G    C    <NA> <NA> <NA>
## [16] <NA> <NA> <NA> <NA> <NA> <NA> D    <NA> A    <NA> <NA> <NA> C
## Levels: A B C D E F G T
```



## 6 MISSING VALUE IMPUTATION

There are some missing values in the dataset. Since, we don't have many observations, we are better not to delete any observations or columns with missing values. Hence, we are trying to replace missing values with a sensible values.

### 6.1 Embarked

Let's take a look at Embarked variable.

```
full[c(62, 830), 'Embarked']
```


```
## [1] <NA> <NA>  
## Levels: C Q S
```

Passengers 62 and 830 are missing Embarkment variable. We can infer their values based on passenger class and fare.

```
embark_fare <- full %>%  
  filter(PassengerId != 62 & PassengerId != 830)
```

Use ggplot2 to visualize Embarked, Class and median Fare variables.

```
ggplot(embark_fare, aes(x = Embarked, y = Fare, fill = factor(Pclass))) +  
  geom_boxplot() +  
  geom_hline(aes(yintercept=80),  
    colour='red', linetype='dashed', lwd=2) +  
  scale_y_continuous(labels=dollar_format())
```



The median fare for a first class passenger departing from Charbourg ('C') matches with the \$80 Fare variable. Hence, we can replace the NA values with 'C'.

```
full$Embarked[c(62, 830)] <- 'C'
full[c(62, 830), 'Embarked']
```

```
## [1] C C
## Levels: C Q S
```

## 6.2 Fare

Similar to Embarked variable, Fare variable on row 1044 has an NA value.

```
full[1044, "Fare"]
```

```
## [1] NA
```


Let's show the all observation.

```
full[1044,]
```

```
##      PassengerId Survived Pclass      Name Sex Age SibSp Parch
## 1044          1044    <NA>      3 Storey, Mr. Thomas male 60.5    0    0
##      Ticket Fare Cabin Embarked Title Fsize  FsizeD Deck
## 1044   3701   NA <NA>      S    Mr    1 singleton <NA>
```

This passenger departed from Southampton ('S'). Let's visualize Fares among same class and embarkment.

```
ggplot(full[full$Pclass == '3' & full$Embarked == 'S', ],
  aes(x = Fare)) +
  geom_density(fill = '#99d6ff', alpha=0.4) +
  geom_vline(aes(xintercept=median(Fare, na.rm=T)),
    colour='red', linetype='dashed', lwd=1) +
  scale_x_continuous(labels=dollar_format())
```



From this visualization, it seems quite reasonable to replace the NA values with median for their class and embarkment.

```
full$Fare[1044] <- median(full[full$Pclass == '3' & full$Embarked == 'S', ]$Fare, na.rm = TRUE)
```

```
full[1044, "Fare"]
```

```
## [1] 8.05
```

### 6.3 Age

There are many missing Age variable in the dataset. Let's look into the number of missing Age variables.

```
sum(is.na(full$Age))
```

```
## [1] 263
```

There are 263 missing Age variables. We can use mice package for guessing missing Age variables.

```
set.seed(129) # Set a random seed
mice_mod <- mice(full[, !names(full) %in% c('PassengerId', 'Name', 'Ticket', 'Cabin', 'Family', 'Surname', 'Survived')])
```

```
##
```

```
## iter imp variable
## 1 1 Age Deck
## 1 2 Age Deck
## 1 3 Age Deck
## 1 4 Age Deck
## 1 5 Age Deck
```


```
## 2 1 Age Deck
## 2 2 Age Deck
## 2 3 Age Deck
## 2 4 Age Deck
## 2 5 Age Deck
## 3 1 Age Deck
## 3 2 Age Deck
## 3 3 Age Deck
## 3 4 Age Deck
## 3 5 Age Deck
## 4 1 Age Deck
## 4 2 Age Deck
## 4 3 Age Deck
## 4 4 Age Deck
## 4 5 Age Deck
## 5 1 Age Deck
## 5 2 Age Deck
## 5 3 Age Deck
## 5 4 Age Deck
## 5 5 Age Deck
```

```
## Warning: Number of logged events: 52
```

```
mice_output <- complete(mice_mod) # Save the output
```

Let's compare the results with the original distribution.

```
par(mfrow=c(1,2))
hist(full$Age, freq=F, main='Original Data', ylim=c(0,0.04))
hist(mice_output$Age, freq=F, main='MICE Output Data', ylim=c(0,0.04))
```



Distributions looks very similar, Let's replace the Age variables with the mice model.

```
full$Age <- mice_output$Age
sum(is.na(full$Age))
```

```
## [1] 0
```

## 7 FEATURE ENGINEERING 2


We can create a couple of new age-dependent variables: Child and Mother.

- A child is someone under 18 years.
- A mother is someone who is female, older than 18, doesn't have the 'Miss' title and has more than 0 children.

### 7.1 Child

Let's look at the relationship between age and survival

```
ggplot(full[1:891,], aes(Age, fill = factor(Survived))) +
  geom_histogram() +
  facet_grid(.~Sex)
```



Let's create a Child variable.

```
full$Child[full$Age < 18] <- 'Child'
full$Child[full$Age >= 18] <- 'Adult'
table(full$Child, full$Survived)
```

```
##
##           0    1
##   Adult 480 275
##   Child  69  67
```

```
full$Child <- factor(full$Child)
```

## 7.2 Mother

Similarly, let's create a Mother variable.

```
full$Mother <- 'Not Mother'
full$Mother[full$Sex == 'female' & full$Parch > 0 & full$Age > 18 & full$Title != 'Miss'] <- 'Mother'
table(full$Mother, full$Survived)
```

```
##
##           0    1
##   Mother      16  38
##   Not Mother 533 304
```

```
full$Mother <- factor(full$Mother)
```

## 8 PREDICTION

We successfully fill the missing values and created new variables in the Titanic datasets. In this section, we will build some models to predict the Survival variable. Let's first check the datasets again using the completeness function.

```
sapply(full, function(x) {sum(is.na(x))})
```

```
## PassengerId   Survived    Pclass      Name      Sex      Age
##           0       418         0         0         0         0
##      SibSp     Parch     Ticket     Fare     Cabin Embarked
##           0         0         0         0      1014         0
##      Title     Fsize   FsizeD     Deck     Child     Mother
##           0         0         0     1014         0         0
```

Let's factorize some variables.

```
full$Sex <- as.factor(full$Sex)
full$Embarked <- as.factor(full$Embarked)
full$Title <- as.factor(full$Title)
full$FsizeD <- as.factor(full$FsizeD)
```

As you can see, we have many new variables. Also, we replaced some missing values with their sensible values. Again, we need to convert some variables into factor.

## 8.1 Split Dataset

Let's split the data back into a training and test datasets.

```
train <- full[1:891,]  
test <- full[892:1309,]
```

## 8.2 RandomForest Model

Let's build the RandomForest model.


```
set.seed(2017)  
random_forrest_model <- randomForest(factor(Survived) ~ Pclass + Sex + Age + SibSp + Parch +  
                                     Fare + Embarked + Title +  
                                     FsizeD + Child + Mother, data = train)
```

Let's predict the Survival column and write output into a file.

```
random_forrest_predictions <- predict(random_forrest_model, test)  
random_forrest_solutions <- data.frame(PassengerID = test$PassengerId, Survived = random_forrest_predictions)  
write.csv(random_forrest_solutions, file = 'output/random_forrest_solutions.csv', row.names = F)
```

The final result for RandomForest model is **0.78947**. Let's visualize the error rate.

```
plot(random_forrest_model, ylim=c(0,0.36))  
legend('topright', colnames(random_forrest_model$err.rate), col=1:3, fill=1:3)
```




The black line shows the overall error rate. The red and green lines show the error rate of 'died' and 'survived'.

### 8.2.1 Visualize the Importance

Let's visualize the importance of variables.

```
random_forrest_importance <- varImp(random_forrest_model, scale = FALSE)
importance_scores <- data.frame(Variables = row.names(random_forrest_importance), MeanDecreaseGini = ra
ggplot(importance_scores, aes(x=reorder(Variables, MeanDecreaseGini), y=MeanDecreaseGini, fill=MeanDecr
  geom_bar(stat='identity') + coord_flip() + theme(legend.position="none") + labs(x="") +
  ggtitle('Variable Importance Random Forest') + theme(plot.title = element_text(hjust = 0.5))
```



It is clear in the plot that, the Title variable has the highest relative importance.

### 8.3 SVM Model